

Elements of Statistical Learning

Junrui Di

Contents

Chapter 2. Overview of Supervised Learning	3
0. Notations	3
1. Types of variables	4
2. Two simple approaches to prediction: least squares and nearest neighbors	4
3. Statistical decision theory	4
4. Function approximation	6
5. Restricted estimators	6
6. Model selection and the Bias-variance tradeoff	6
Chapter 3: Linear Methods for Regression	8
1. Introduction	8
2. Linear regression models and least square	8
3. Subset selection	9
4. Shrinkage methods	9
5. Methods using derived input directions	11
Chapter 4: Linear Methods for Classification	13
1. Linear discriminant analysis	13
2. Logistic regression	13
3. Separating hyperplanes	15
Chapter 5: Basis Expansions and Regularization	16
1. Introduction	16
2. Piecewise Polynomials and Splines (restricted model)	16
3. Filtering and Feature Extraction	16
4. Smoothing Splines	16
4.1 Degree of Freedom and Smoother Matrices	17
5. Automatic Selection of Smoothing Parameters	17
6. Nonparametric Logistic Regression	17

7. Multidimensional Splines	18
8. Regularization and Reproducing Kernel Hilbert Spaces	18
9. Wavelet Smoothing	18
Chapter 6: Kernel Smoothing Method	19
0. Introduction	19
1. One-dimensional Kernel Smoothers	19
2. Selecting Width of the Kernel	19
3. Kernel Density Estimation and Classification	19
3.1 Naive Bayes	19
4. Radial Basis Functions and Kernels	19
Chapter 7: Model Assessment and Selection	20
1. Introduction	20
2. Bias, Variance, and MOdel Complexity	20
3. Bia-Variance Decomposition	20
4. The Optimism of the Training Error Rate	20
5. Estimates of In-Sample Prediction Error	20
6. Effective Number of Parameters	20
7. The Bayesian Approachj and BIC	20
10. Cross Validation	21
11. Bootstrap	21
Chapter 8: Model Inference and Averaging	22
1. Introduction	22
2. The Bootstrap and Maximum Likelihood Methods	22
3. Baysian Methods	22
4. Boostrape v.s Bayesian Inference	23
5. The EM Algorithm	23
6. MCMC for Sampling from the Posterior	23
7. Bagging	23
8. Model Averaging and Stacking	23
Chapter 9: Additive Models, Trees, and Related Methods	24
1. Generalized Additive Model	24
2. Tree-Based Methods	24
3. PRIM: Bump Hunting	26
4. MARS:; Multivariate Adaptive Regression Splines	26
5. Hierachical Mixtures of Experts	26
6. Missing Data	28

Chapter 10: Boosting and Additive Trees	29
1. Introduction	29
Chapter 11: Neural Networks	30
1. Introduction	30
Chapter 12: Support Vector Machines and Flexible Discriminants	31
1. Introduction	31
Chapter 13: Prototype Methods and Nearest-Neighbors	32
1. Introduction	32
Chapter 14: Unsupervised Learning	33
1. Introduction	33
Chapter 15: Random Forests	34
1. Introduction	34
Chapter 16: Ensemble Learning	35
1. Introduction	35
Chapter 17: Undirected Graphical Models	36
1. Introduction	36
Chapter 18: High-Dimensional Problems: $p \gg N$	37
1. Introduction	37

Chapter 2. Overview of Supervised Learning

0. Notations

- Use upper case letters X, Y, G for generic variables
 - Input variable X with j th component denoted as X_j
 - Quantitative output Y
 - Qualitative output G
- Observed values in lowercase
 - i th observation of X is x_i (a scalar or a vector)
- Matrices are represented in bold uppercase letters
 - A set of N input p -vectors $x_i, i = 1 \dots N$ will be $\mathbf{X} \in \mathbb{R}^{N \times p}$
 - p -vector of input x_i for the i th observation v.s. the N -vector \mathbf{x}_j for all the observations on variable X_j
 - All vectors are assumed to be column vectors, the i th row of \mathbf{X} is x_i^T .

1. Types of variables

- Qualitative variables, factors, categorical or discrete variables → **Classification**
- Quantitative measurements → **Regression**
- Ordered qualitative variables

2. Two simple approaches to prediction: least squares and nearest neighbors

2.1 Linear models and least squares

- Linear model
- Input: $X^T = (X_1, X_2, \dots, X_p)$, Outcome: Y
- Model: $\hat{Y} = \hat{\beta}_0 + \sum_{j=1}^p X_j \hat{\beta}_j$ or $\hat{Y} = X^T \hat{\beta}$
- Least Square
- To minimize $\text{RSS}(\beta) = \sum_{i=1}^N (y_i - x_i^T \beta)^2$ or $\text{RSS}(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$
- Differentiate w.r.t. β gives $\mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta) = 0$
- Solves to $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- Fitted value at the i th input x_i is $\hat{y}_i = x_i^T \hat{\beta}$

2.2 Nearest neighbor methods

The k -nearest neighbor fit for \hat{Y} : $\hat{Y}(x) = \frac{1}{k} \sum_{x_i \in N_k(x)} y_i$, where $N_k(x)$ is the neighborhood of x as the k closest points x_i in the training set.

For k -nearest neighbor fit, the error on the training data should be approximately an increasing function of k , and 0 for $k = 1$. We cannot use sum of squared errors as training criterion for picking k .

There is only one parameter in the fit, which is k . But the effective number of parameters is N/k , because there would be N/k neighbors and we need that many means for each of the neighborhood.

2.3 From least square to nearest neighbors

Least square: smooth linear decision boundary and stable to fit, but heavily rely on assumption of linear decision boundary. **Low variance but high bias.**

knn : no strong assumption and can adapt to any situation, but unstable (depend on a handful of input points and their positions). **High variance but low bias.**

3 Statistical decision theory

[1.] *Quantitative output* framework:

- Output $Y \in \mathbb{R}$, and input $X \in \mathbb{R}^p$
- Joint distribution $\Pr(X, Y)$
- Goal: find a function $f(X)$ to predict Y .

- Loss $L(Y, f(X))$, e.g. a squared error loss $L(Y, f(X))$

The expected squared prediction value is

$$\begin{aligned} EPE(f) &= E(Y - f(X))^2 \\ &= E_X E_{Y|X}([Y - f(X)]^2 | X) \quad \text{conditioning on } X \end{aligned}$$

which can be minimized by $f(x) = \operatorname{argmin}_c E_{Y|X}([Y - c]^2 | X)$, which can be solved by $f(x) = E(Y | X = x)$, also known as the regression function. **The best prediction of Y at any point $X = x$ is the conditional mean when best is measured by average squared error.**

knn mimics this framework, by $\hat{f}(x) = \operatorname{Ave}(y_i | x_i \in N_k(x))$ with two approximations

- expectation is approximated by averaging over sample data;
- conditioning at a point is relaxed to conditioning on some region close to the target point

Least square also mimics this framework, with the assumption that the regression function $f(x)$ is approximately linear in its argument, i.e. $f(x) \approx x^T \beta$. Therefore, β can be solved by $\beta = [E(XX^T)]^{-1} E(XY)$. That is not to condition on X , rather we used our knowledge of the functional relationship to pool over values of X . Least square estimates replace $E(\cdot)$ by averaging over the training data.

[2.] *Qualitative output framework:*

- Suppose there are K classes in G .
- Loss function can be represented by a $K \times K$ matrix \mathbf{L} , where each position $L(k, l)$ is the loss for misclassifying G_k as G_l . Most commonly we can use the zero_one loss, that is the set the the loss as 1.

The expected prediction error is

$$\begin{aligned} EPE &= E[L(G, \hat{G}(X))] \\ &= E_X \sum_{k=1}^K L[G_k, \hat{G}(X)] Pr(G_k | X) \end{aligned}$$

which can be minimized by

$$\begin{aligned} \hat{G}(x) &= \operatorname{argmin}_{g \in G} \sum_{k=1}^K L[G_k, g] Pr(G_k | X = x) \\ &= \operatorname{argmin}_{g \in G} [1 - Pr(g | X = x)] \\ &= \max Pr(g | X = x) \end{aligned}$$

This is known as the *Bayes Classifier*, such that we classify to the most probably class, using the conditional distribution $Pr(G|X)$.

knn directly approximates this solution using majority vote in a nearest neighborhood, except that conditional probability at a point is relaxed to conditional probability within a neighborhood of a point and probability are approximated by training sample proportions.

4. Function approximation

Data $\{x_i, y_i\}$ are considered to be from a $p + 1$ dimensional Euclidean space. The function $f(x)$ has domain equal to a p -dimensional subspace. Data and function are related via the model $y_i = f(x_i) + \epsilon_i$. The goal for learning is to find an approximation to $f(x)$ in \mathbb{R}^p given the representation in the domain of data which is \mathbb{R}^{p+1}

The question is to find a set of parameters θ for the function $f_\theta(x)$ with following criterion

[1.] *Least square*

To minimize the $RSS(\theta) = \sum_{i=1}^N (y_i - f_\theta(x_i))^2$

[2.] *Maximum likelihood estimation*

If we have a random sample $y_i, i = 1 \dots N$ from a density $Pr_\theta(y)$. The log-probability of the observed sample is $L(\theta) = \sum_i \log Pr_\theta(y_i)$. The most reasonable values for θ are those for which the probability of the observed sample is the largest.

5. Restricted estimators

Minimizing the RSS leads to many solution, because any function \hat{f} that passing through the training points is a solution. Therefore, we need to add complexity restrictions, that is, for all input points x sufficiently close to each other in some metric, \hat{f} exhibits some special structure such as nearly constant, linear, or low-order polynomial behavior.

5.1 Roughness penalty and Bayesian methods

$$PRSS(f; \lambda) = RSS(f) + \lambda J(f)$$

with penalty $J(\cdot)$. E.g. cubic smoothing splines penalizes on large values of second order derivative.

5.2 Kernel methods and local regression

Kernel methods control the nature of the local neighborhood, using a kernel function $K_\lambda(x_0, x)$, which put weights to points x in a region near x_0 (λ controls the width of the neighborhood).

A local regression estimate of $f(x_0)$ as $f_{\hat{\theta}}(x_0)$ where $\hat{\theta}$ minimizes $RSS(f_\theta, 0) = \sum_i K_\lambda(x_0, x_i)(y_i - f_\theta(x_i))^2$.

5.3 Basis functions and dictionary methods

$$f_\theta(x) = \sum_{m=1}^M \theta_m h_m(x)$$

6. Model selection and the Bias-variance tradeoff

Data $\{x_i, y_i\}$, model $y = f(x) + \epsilon$, where $E(\epsilon) = 0$, $var(\epsilon) = \sigma^2$

$$E[(y - \hat{f}(x))^2] = (\text{Bias}[\hat{f}(x)])^2 + \text{Var}[\hat{f}(x)] + \sigma^2$$

* $\text{Bias}[\hat{f}(x)] = E[\hat{f}(x)] - f(x)$: error caused by simplifying assumptions build into the method

- $\text{Var}[\hat{f}(x)] = E[(E[\hat{f}(x)] - \hat{f}(x))^2]$: variance of the learning method
- irreducible error σ^2 due to the new test target.

Derivation




Figure 1: Model Complexity v.s. Prediction Errors

Derivation [\[edit \]](#)

The derivation of the bias–variance decomposition for squared error proceeds as follows.^{[9][10]} For notational convenience, we abbreviate $f = f(x)$, $\hat{f} = \hat{f}(x; D)$ and we drop the D subscript on our expectation operator. recall that, by definition, for any random variable X , we have

$$\text{Var}[X] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2.$$

Rearranging, we get:

$$\mathbb{E}[X^2] = \text{Var}[X] + (\mathbb{E}[X])^2.$$

Since f is **deterministic**, i.e. independent of D ,

$$\mathbb{E}[f] = f.$$

Thus, given $y = f + \varepsilon$ and $\mathbb{E}[\varepsilon] = 0$ (because ε is noise), implies $\mathbb{E}[y] = \mathbb{E}[f + \varepsilon] = \mathbb{E}[f] = f$.

Also, since $\text{Var}[\varepsilon] = \sigma^2$,

$$\text{Var}[y] = \mathbb{E}[(y - \mathbb{E}[y])^2] = \mathbb{E}[(y - f)^2] = \mathbb{E}[(f + \varepsilon - f)^2] = \mathbb{E}[\varepsilon^2] = \text{Var}[\varepsilon] + (\mathbb{E}[\varepsilon])^2 = \sigma^2 + 0^2 = \sigma^2.$$

Thus, since ε and \hat{f} are independent, we can write

$$\begin{aligned} \mathbb{E}[(y - \hat{f})^2] &= \mathbb{E}[(f + \varepsilon - \hat{f})^2] \\ &= \mathbb{E}[(f + \varepsilon - \hat{f} + \mathbb{E}[\hat{f}] - \mathbb{E}[\hat{f}])^2] \\ &= \mathbb{E}[(f - \mathbb{E}[\hat{f}])^2] + \mathbb{E}[\varepsilon^2] + \mathbb{E}[(\mathbb{E}[\hat{f}] - \hat{f})^2] + 2\mathbb{E}[(f - \mathbb{E}[\hat{f}])\varepsilon] + 2\mathbb{E}[\varepsilon(\mathbb{E}[\hat{f}] - \hat{f})] + 2\mathbb{E}[(\mathbb{E}[\hat{f}] - \hat{f})(f - \mathbb{E}[\hat{f}])] \\ &= (f - \mathbb{E}[\hat{f}])^2 + \mathbb{E}[\varepsilon^2] + \mathbb{E}[(\mathbb{E}[\hat{f}] - \hat{f})^2] + 2(f - \mathbb{E}[\hat{f}])\mathbb{E}[\varepsilon] + 2\mathbb{E}[\varepsilon]\mathbb{E}[\mathbb{E}[\hat{f}] - \hat{f}] + 2\mathbb{E}[\mathbb{E}[\hat{f}] - \hat{f}](f - \mathbb{E}[\hat{f}]) \\ &= (f - \mathbb{E}[\hat{f}])^2 + \mathbb{E}[\varepsilon^2] + \mathbb{E}[(\mathbb{E}[\hat{f}] - \hat{f})^2] \\ &= (f - \mathbb{E}[\hat{f}])^2 + \text{Var}[\varepsilon] + \text{Var}[\hat{f}] \\ &= \text{Bias}[\hat{f}]^2 + \text{Var}[\varepsilon] + \text{Var}[\hat{f}] \\ &= \text{Bias}[\hat{f}]^2 + \sigma^2 + \text{Var}[\hat{f}]. \end{aligned}$$

Finally, MSE loss function (or negative log-likelihood) is obtained by taking the expectation value over $x \sim P$:

$$\text{MSE} = \mathbb{E}_x \left\{ \text{Bias}_D[\hat{f}(x; D)]^2 + \text{Var}_D[\hat{f}(x; D)] \right\} + \sigma^2.$$

Figure 2: Bias-Variance Tradeoff

Chapter 3: Linear Methods for Regression

1. Introduction

Linear regression assumes that the regression function $E(Y|X)$ is linear in the inputs X_1, \dots, X_p .

2. Linear regression models and least square

[1.] *Linear regression from a least square point of view* (minimal assumption about the distribution)

- Form: $f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$
- Data: $\{x_i, y_i\}$ $i = 1 \dots N$, each $x_i = (x_{i1} \dots x_{ip})^T$ is a feature vector, with parameters $\beta = (\beta_0, \beta_1, \dots, \beta_p)^T$
- Least square: To minimize $RSS(\beta) = \sum_{i=1}^N (y_i - f(x_i))^2 = \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2$ or $RSS(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T(\mathbf{y} - \mathbf{X}\beta)$ in matrix form. **LSE makes no assumptions about the validity of the model form**
- LSE: $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- Fitted value: $\hat{\mathbf{y}} = \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$. $\mathbf{H} = \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}$ is the projector of \mathbf{y} onto the subspace spanned by column space of \mathbf{X} .
- Inference on parameters (assuming y_i 's are uncorrelated and gave constant variance σ^2 , and x_i are fixed)
 - $Var(\hat{\beta}) = (\mathbf{X}^T \mathbf{X})^{-1} \sigma^2$
 - $\hat{\sigma}^2 = \frac{1}{N-p-1} \sum_{i=1}^N (y_i - \hat{y}_i)^2$

[2.] *Linear regression with Gaussian error*

- Model Assumption: $Y = \beta_0 + \sum_{j=1}^p X_j \beta_j + \epsilon$, where $\epsilon \sim N(0, \sigma^2)$
- Distributional properties of model parameters
 - $\hat{\beta} \sim N(\beta, (\mathbf{X}^T \mathbf{X})^{-1} \sigma^2)$
 - $(N - p - 1) \hat{\sigma}^2 \sim \sigma^2 \chi_{N-p-1}^2$
 - $\hat{\beta}$ and $\hat{\sigma}^2$ are statistically independent.
- Inference on single parameter β_j

Under $H_o : \beta_j = 0$, $z_j = \frac{\hat{\beta}_j}{\hat{\sigma} \sqrt{(\mathbf{X}^T \mathbf{X})_{jj}^{-1}}} \sim t_{N-p-1}$, and β_j has a $1 - 2\alpha$ confidence interval of $(\hat{\beta}_j - z^{1-\alpha} \hat{\sigma} \sqrt{(\mathbf{X}^T \mathbf{X})_{jj}^{-1}}, \hat{\beta}_j + z^{1-\alpha} \hat{\sigma} \sqrt{(\mathbf{X}^T \mathbf{X})_{jj}^{-1}})$

- Nested Model Comparison (test whether the added variables are necessary to the model)

$$F = \frac{(RSS_0 - RSS_1)/(p_1 - p_0)}{RSS_1/(N - p_1 - 1)} \sim F_{p_1 - p_0, N - p_1 - 1}, \quad \text{where } RSS_1 \text{ is for the larger model}$$

2.1 The Gauss-Markow Theorem

Least square estimates of β have the smallest variance among all linear unbiased estimates.

The least square estimator to estimate parameters $\theta = \alpha^T \beta$ is $\hat{\theta} = \alpha^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$. It is an unbiased estimator, i.e. $E(\alpha^T \hat{\beta}) = \alpha^T \beta$. Gauss-Markow theorem states that $Var(\alpha \hat{\beta})$ has the smallest variance for any unbiased estimator.

We may want to trade a little bias for larger reduction in variance.

2.2 Regression by successive orthogonalization

Algorithm 3.1 *Regression by Successive Orthogonalization.*

1. Initialize $\mathbf{z}_0 = \mathbf{x}_0 = \mathbf{1}$.
 2. For $j = 1, 2, \dots, p$
 - Regress \mathbf{x}_j on $\mathbf{z}_0, \mathbf{z}_1, \dots, \mathbf{z}_{j-1}$ to produce coefficients $\hat{\gamma}_{\ell j} = \langle \mathbf{z}_\ell, \mathbf{x}_j \rangle / \langle \mathbf{z}_\ell, \mathbf{z}_\ell \rangle$, $\ell = 0, \dots, j-1$ and residual vector $\mathbf{z}_j = \mathbf{x}_j - \sum_{k=0}^{j-1} \hat{\gamma}_{kj} \mathbf{z}_k$.
 3. Regress \mathbf{y} on the residual \mathbf{z}_p to give the estimate $\hat{\beta}_p$.
-

Figure 3: Gram-Schmidt procedure for multiple regression

2.3 Multiple outcomes

Data: $Y_1 \dots Y_K$, with the model $Y_k = \beta_{0k} + \sum_{j=1}^p X_j \beta_{jk} + \epsilon_k$, with the matrix form $\mathbf{Y} = \mathbf{XB} + \mathbf{E}$, where \mathbf{Y} is $N \times K$, \mathbf{X} is $N \times p + 1$, and \mathbf{B} is $(p + 1) \times K$.

$RSS(\mathbf{B}) = \sum_k \sum_i (y_{ik} - f_k(x_i))^2 = \text{tr}(\mathbf{Y} - \mathbf{XB})^T (\mathbf{Y} - \mathbf{XB})$ is the RSS with LSE $\hat{\mathbf{B}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$

3. Subset selection

- a. Best subset selection
- b. Forward and backward selection

4. Shrinkage methods

4.1 Ridge regression

- RSS

$$\hat{\beta}^{\text{ridge}} = \underset{\beta}{\text{argmin}} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

or in the matrix form

$$\text{RSS}(\lambda) = (\mathbf{y} - \mathbf{X}\beta)^T(\mathbf{y} - \mathbf{X}\beta) + \lambda\beta^T\beta$$

with the solution

$$\hat{\beta}^{\text{ridge}} = (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}^T\mathbf{y}$$

Even if $\mathbf{X}^T\mathbf{X}$ is not of full rank, $(\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})$ is still nonsingular.

- Degree of freedom $\text{df}(\lambda) = \text{tr}[\mathbf{X}(\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})^{-1}\mathbf{X}^T] = \sum_{j=1}^p \frac{d_j^2}{d_j^2 + \lambda}$
- Ridge solutions are not equivariant under scaling of the inputs, and one normally standardizes the inputs before solving for estimation.

4.2 Lasso

- RSS

$$\hat{\beta}^{\text{lasso}} = \underset{\beta}{\text{argmin}} \left\{ \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^p x_{ij}\beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\}$$

* Shrinkage $s = t / \sum_j |\hat{\beta}_j|$ where $\hat{\beta}_j$ is the least square estimation.

4.3 Subset selection, ridge, and lasso

[1.] *Orthonormal input matrix* \mathbf{X}

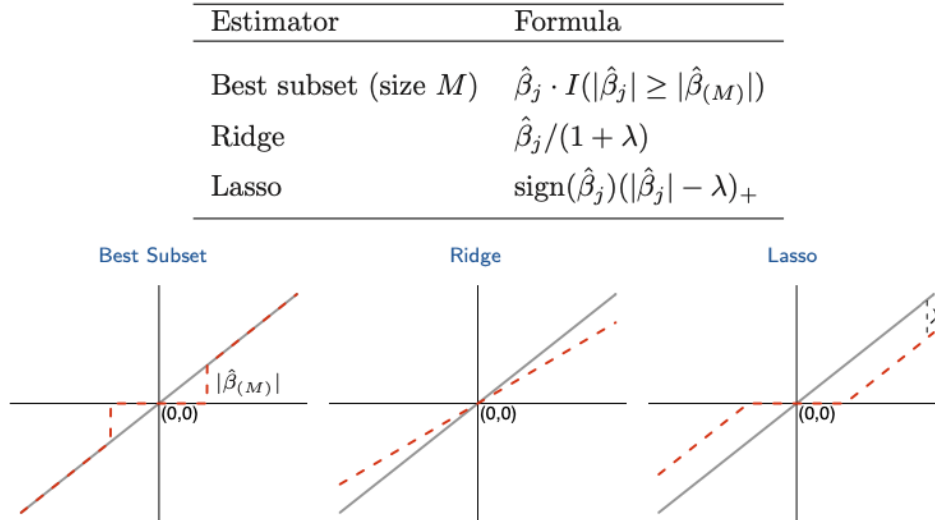


Figure 4: Gram-Schmidt procedure for multiple regression

- Ridge: proportional shrinkage
- LASSO: translate by a constant factor and truncating at zero, i.e soft thresholding
- Best subject: drops all the variables with coefficient smaller than the M th largest, i.e. hard thresholding

[2.] *Nonorthogonal case*




Figure 5: Gram-Schmidt procedure for multiple regression

Elastic net

$$\lambda \sum_{j=1}^p (\alpha \beta_j^2 + (1 - \alpha) |\beta_j|)$$

4.4 Least angle regression

5. Methods using derived input directions

5.1 Principal components regression

PC regression forms the derived input columns $\mathbf{z}_m = \mathbf{X}v_m$ and then regresses \mathbf{y} on $\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_M$. Since they are orthogonal, each parameter is simply $\hat{\theta}_n = \frac{\langle \mathbf{z}_m, \mathbf{y} \rangle}{\langle \mathbf{z}_m, \mathbf{z}_m \rangle}$. It can be converted back to $\hat{\beta}_M^{pcr} = \sum_{m=1}^M \hat{\theta}_m v_m$.

The m th principal component direction v_m solveS:

$$\begin{aligned} & \max_{\alpha} \text{Var}(\mathbf{X}\alpha) \\ & \text{subject to } \|\alpha\| = 1, \alpha^T \mathbf{S}v_l = 0, \quad l = 1, \dots, n-1 \end{aligned}$$

where \mathbf{S} is the sample covariance

Algorithm 3.3 *Partial Least Squares.*

1. Standardize each \mathbf{x}_j to have mean zero and variance one. Set $\hat{\mathbf{y}}^{(0)} = \bar{y}\mathbf{1}$, and $\mathbf{x}_j^{(0)} = \mathbf{x}_j$, $j = 1, \dots, p$.
 2. For $m = 1, 2, \dots, p$
 - (a) $\mathbf{z}_m = \sum_{j=1}^p \hat{\varphi}_{mj} \mathbf{x}_j^{(m-1)}$, where $\hat{\varphi}_{mj} = \langle \mathbf{x}_j^{(m-1)}, \mathbf{y} \rangle$.
 - (b) $\hat{\theta}_m = \langle \mathbf{z}_m, \mathbf{y} \rangle / \langle \mathbf{z}_m, \mathbf{z}_m \rangle$.
 - (c) $\hat{\mathbf{y}}^{(m)} = \hat{\mathbf{y}}^{(m-1)} + \hat{\theta}_m \mathbf{z}_m$.
 - (d) Orthogonalize each $\mathbf{x}_j^{(m-1)}$ with respect to \mathbf{z}_m : $\mathbf{x}_j^{(m)} = \mathbf{x}_j^{(m-1)} - [\langle \mathbf{z}_m, \mathbf{x}_j^{(m-1)} \rangle / \langle \mathbf{z}_m, \mathbf{z}_m \rangle] \mathbf{z}_m$, $j = 1, 2, \dots, p$.
 3. Output the sequence of fitted vectors $\{\hat{\mathbf{y}}^{(m)}\}_1^p$. Since the $\{\mathbf{z}_\ell\}_1^m$ are linear in the original \mathbf{x}_j , so is $\hat{\mathbf{y}}^{(m)} = \mathbf{X} \hat{\beta}^{\text{pls}}(m)$. These linear coefficients can be recovered from the sequence of PLS transformations.
-

Figure 6: Gram-Schmidt procedure for multiple regression

5.2 Partial least square

The m th PLS direction $\hat{\psi}_m$ solves:

$$\begin{aligned} & \max_{\alpha} \text{Corr}^2(\mathbf{y}, \mathbf{X}\alpha) \text{Var}(\mathbf{X}\alpha) \\ & \text{subject to } \|\alpha\| = 1, \alpha^T \hat{\mathbf{S}} \hat{\psi}_l = 0, \quad l = 1, \dots, m-1 \end{aligned}$$

Chapter 4: Linear Methods for Classification

1. Linear discriminant analysis

Question set up for classification:

- Goal: To know the class posteriors $Pr(G|X)$ for optimal classification
- Parameters: $f_k(x)$ is the class conditional density of X in class $G = k$, and π_k is the probability of class k , with $\sum_{k=1}^K \pi_k = 1$.
- Bayes theorem gives that $Pr(G = k|X = x) = \frac{f_k(x)\pi_k}{\sum_{l=1}^K f_l(x)\pi_l}$

Suppose each class density is a multivariate Gaussian

$$f_k(x) = \frac{1}{(2\pi)^p |\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1}(x-\mu_k)} \quad \text{assume equal variance across classes}$$

To compare the two classes k and l , we have

$$\log \frac{Pr(G = k|X = x)}{Pr(G = l|X = x)} = \log \frac{\pi_k}{\pi_l} - \frac{1}{2}(\mu_k + \mu_l)^T \Sigma^{-1}(\mu_k - \mu_l) + x^T \Sigma^{-1}(\mu_k - \mu_l)$$

which is linear in x .

- Linear discriminant function is to solve for $G(x) = \operatorname{argmax}_k \delta_k(x)$

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$$

π_k, μ_k, Σ all needs to be estimated empirically.

- Quadratic discriminant function is where Σ_k are not all the same, then the function becomes

$$\delta_k(x) = -\frac{1}{2} \log |\Sigma_k| + \log \pi_k - \frac{1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k)$$

1.1 Regularized discriminant analysis

Allow one to shrink the separate covariance of QDA toward a common covariance as in LDA

$$\hat{\Sigma}_k(\alpha) = \alpha \hat{\Sigma}_k + (1 - \alpha) \hat{\Sigma}$$

2. Logistic regression

For multiple K class,

$$Pr(G = k|X = x) = \frac{\exp(\beta_{k0} + \beta_k^T x)}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)} \quad k = 1, \dots, K-1$$
$$Pr(G = l|X = x) = \frac{1}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)}$$

2.1 Fitting a binary class logistic regression

- Log likelihood for multiclass

$$l(\theta) = \sum_{i=1}^N \log p_{g_i}(x_i; \theta)$$

where $p_k(x_i; \theta) = \Pr(G = k | X = x_i; \theta)$.

- For a two class case, $p_1(x; \theta) = p(x; \theta)$ corresponding to $y_i = 1$, and $p_2(x; \theta) = 1 - p(x; \theta)$ corresponding to $y_i = 0$. Then the log likelihood becomes

$$\begin{aligned} l(\beta) &= \sum_{i=1}^N \{y_i \log p(x_i; \theta) + (1 - y_i) \log(1 - p(x_i; \theta))\} \\ &= \sum_i \{y_i \beta^T x_i - \log(1 + e^{\beta^T x_i})\} \end{aligned}$$

- First order derivative

$$\begin{aligned} \frac{\partial l(\beta)}{\partial \beta} &= \sum_i x_i (y_i - p(x_i; \beta)) = 0 \\ &= \mathbf{X}^T (\mathbf{y} - \mathbf{p}) \end{aligned}$$

- Second order derivative

$$\begin{aligned} \frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T} &= - \sum_i x_i x_i^T p(x_i; \beta) (1 - p(x_i; \beta)) \\ &= -\mathbf{X}^T \mathbf{W} \mathbf{X} \end{aligned}$$

- Newton Raphson

$$\begin{aligned} \beta^{\text{new}} &= \beta^{\text{old}} - \left(\frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T} \right)^{-1} \frac{\partial l(\beta)}{\partial \beta} \\ &= \beta^{\text{old}} + (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{p}) \\ &= (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} (\mathbf{X} \beta_{\text{old}} + \mathbf{W}^{-1} (\mathbf{y} - \mathbf{p})) \\ &= (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} \mathbf{z} \end{aligned}$$

The last step can be considered as iteratively reweighted least squares

$$\beta^{\text{new}} \rightarrow \operatorname{argmin}_{\beta} (\mathbf{z} - \mathbf{X}\beta)^T \mathbf{W} (\mathbf{z} - \mathbf{X}\beta)$$

2.1 L1 regularized logistic regression

$$\max_{\beta_0, \beta_1} \left\{ \sum_{i=1}^N [y_i (\beta_0 + \beta^T x_i) - \log(1 + e^{\beta_0 + \beta^T x_i})] - \lambda \sum_{j=1}^p |\beta_j| \right\}$$

3. Separating hyperplanes

Hyperplane (affine set) L defined by the equation $f(x) = \beta_0 + \beta^T x = 0$, in \mathbb{R}^2 , is a line, with the properties

- For any two points in L , $\beta^T(x_1 - x_2) = 0$
- For any point x_0 in L , $\beta^T x_0 = -\beta_0$
- The signed distance of any point x to L is $\frac{1}{\|\beta\|}(\beta^T x + \beta_0) = \frac{1}{\|f'(x)\|}f(x)$

3.1 Perceptron learning algorithm

For a two class problem $y_i \in \{-1, 1\}$

$$\begin{aligned} x_i^T \beta + \beta_0 &< 0 & \text{if } y_i = 1 \text{ is misclassified} \\ x_i^T \beta + \beta_0 &> 0 & \text{if } y_i = -1 \text{ is misclassified} \end{aligned}$$

Therefore, the goal is to minimize

$$D(\beta, \beta_0) = - \sum_{i \in \mathcal{M}} y_i (x_i^T \beta + \beta_0)$$

where \mathcal{M} is the set of misclassified points. This quantity is nonnegative and proportional to the distance of the misclassified points to the decision boundary $\beta^T x + \beta_0 = 0$. The gradient is

$$\begin{aligned} \frac{\partial D(\beta, \beta_0)}{\partial \beta} &= - \sum_{i \in \mathcal{M}} y_i x_i \\ \frac{\partial D(\beta, \beta_0)}{\partial \beta_0} &= - \sum_{i \in \mathcal{M}} y_i \end{aligned}$$

3.2 Optimal separating hyperplanes

Definition: OSH separates the two classes and maximizes the distance to the closest point from either class.

$$\begin{aligned} &\max_{\beta, \beta_0, \|\beta\|=1} M \\ &\text{subject to } y_i(x_i^T \beta + \beta_0) \geq M \quad \forall i \end{aligned}$$

Interpretation: all points are at least a signed distance M from the decision boundary defined by β and β_0 , and seek the largest the M . $\|\beta\| = 1$ can be removed by changing the condition to $y_i(x_i^T \beta + \beta_0) \geq M\|\beta\|$

If we arbitrarily set $\|\beta\| = 1/M$, the question becomes

$$\begin{aligned} &\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2 \\ &\text{subject to } y_i(x_i^T \beta + \beta_0) \geq 1 \quad \forall i \end{aligned}$$

The constraints define a margin around the linear decision boundary of thickness $1/\|\beta\|$.

The question is to minimize the Lagrange function

$$L_p = \frac{1}{2} \|\beta\|^2 - \sum_i \alpha_i [y_i(x_i^T \beta + \beta_0) - 1]$$

Chapter 5: Basis Expansions and Regularization

1. Introduction

Core concept: To augment and replace the vector of inputs X with additional variables which are transformations of X and then use the linear models in this new space of derived input features.

$$f(X) = \sum_{m=1}^M \beta_m h_m(X)$$

where $h_m(X) : \mathbb{R}^p \rightarrow \mathbb{R}$ is a transformation of X . This is called a linear basis expansion in X .

2. Peicewise Polynomials and Splines (restricted model)

Dividing the domain of X into contiguous intervals, and representing f by a separate polynomial in each interval.

An order- M (degree of the polynomial plus 1, i.e. cubic spline has $M = 4$) spline with knots $\xi_j, j = 1 \dots K$ is a piecewise polynomial of order M , and has continuous derivatives up to order $M - 2$. The form of the truncate power basis is

$$\begin{aligned} h_j(X) &= X^{j-1}, j = 1 \dots M \\ h_{M+l}(x) &= (X - \xi_l)_+^{M-1}, l = 1 \dots K \end{aligned}$$

Cubic spline is typically good enough to depict continuity unless we need smooth derivatives.

One approach is to parametrize a family of spline by the number of basis functions or degree of freedom and have the observations x_i determine the positions of the knots

2.1 Natural Cubic Splines

An NCS adds additional constraints, namely that the function is linear beyond the boundary knots which frees up four degrees of freedom (two constraints each in both boundary regions)

An NCS with K knots is represented by K basis functions.

$$\begin{aligned} N_1(X) &= 1, N_2(X) = X, N_{k+2}(X) = d_k(X) - d_{k-1}(X) \\ d_k(X) &= \frac{(X - \xi_k)_+^3 - (X - \xi_K)_+^3}{\xi_K - \xi_k} \end{aligned}$$

3. Filtering and Feature Extraction

4. Smoothing Splines

Smoothing splines avoids the knot selection problem completely by using a maximal set of knots. The complexity of the fit is controlled by regularization.

$$RSS(f, \lambda) = \sum_{i=1}^N \{y_i - f(x_i)\}^2 + \lambda \int \{f''(t)\}^2 dt$$

The penalty term penalizes curvature in the function using the smoothing parameter λ . $\lambda = 0$: f is any function that interpolates the data. $\lambda = \infty$: simple least squares linear fit with no second derivative tolerated. The results has a unique minimizer which is a natural cubic spline with knots at the unique values of the x_i . The solution can be written as

$$f(x) = \sum_{j=1}^N N_j(x) \theta_j$$

where $N_j(x)$ are an N -dimensional set of basis functions for representing this family of natural splines.

The criterion thus reduces to

$$RSS(\theta, \lambda) = (\mathbf{y} - \mathbf{N}\theta)^T (\mathbf{y} - \mathbf{N}\theta) + \lambda \theta^T \mathbf{\Omega}_N \theta$$

where $\{\mathbf{N}\}_{ij} = N_j(x_i)$ and $\{\mathbf{\Omega}_N\}_{jk} = \int N_j''(t) N_k''(t) dt$, with the solution

$$\hat{\theta} = (\mathbf{N}^T \mathbf{B} + \lambda \mathbf{\Omega}_N)^{-1} \mathbf{N}^T \mathbf{y}$$

4.1 Degree of Freedom and Smoother Matrices

Fitted value can be written as

$$\hat{\mathbf{f}} = \mathbf{N}(\mathbf{N}^T \mathbf{B} + \lambda \mathbf{\Omega}_N)^{-1} \mathbf{N}^T \mathbf{y} = \mathbf{S}_\lambda \mathbf{y}$$

\mathbf{S}_λ is known as the smoother matrix, which depends only on x_i and λ .

\mathbf{S}_λ is symmetric and positive semidefinite and has real eigendecomposition.

The effective degree of freedom of a smoothing spline is

$$\text{df}_\lambda = \text{trace}(\mathbf{S}_\lambda)$$

5. Automatic Selection of Smoothing Parameters

Fix the df to determine λ

```
smooth.spline(x,y,df = 6)
```

Covariance: $\text{Cov}(\hat{\mathbf{f}}) = \mathbf{S}_\lambda \text{Cov}(\mathbf{y}) \mathbf{S}_\lambda^T = \mathbf{S}_\lambda \mathbf{S}_\lambda^T$ assuming iid normal error

Bias: $\text{Bias}(\hat{\mathbf{f}}) = \mathbf{f} - \mathbf{S}_\lambda \mathbf{f}$

small df_λ : underfit and trims down the hills and fills in the valleys

large df_λ : overfit and wiggly

6. Nonparametric Logistic Regression

$$\log \frac{\Pr(Y = 1|X = x)}{\Pr(Y = 0|X = x)} = f(x)$$

where $f(x)$ is a smooth function. $\Pr(Y = 1|X = x)$ is therefore also smooth.

The log likelihood is

$$l(f; \lambda) = \sum_{i=1}^N [y_i f(x_i) - \log(1 + e^{f(x_i)})] - \frac{1}{2} \int \{f''(t)\}^2 dt$$

7. Multidimensional Splines

Generalized form

$$\min_f \sum_{i=1}^N B\{y_i - f(x_i)\}^2 + \lambda J[f]$$

where J is an appropriate penalty functional for stabilizing a function f in \mathbb{R}^d . E.g. when $d = 2$, $J[f] = \int_{\mathbb{R}} [(\frac{\partial^2 f(x)}{\partial x_1^2})^2 + 2(\frac{\partial^2 f(x)}{\partial x_1 \partial x_2}) + (\frac{\partial^2 f(x)}{\partial x_2^2})^2] dx_1 dx_2$

8. Regularization and Reproducing Kernel Hilbert Spaces

9. Wavelet Smoothing

Chapter 6: Kernel Smoothing Method

0. Introduction

Class of regression techniques that achieve flexibility in estimating the regression function $f(X)$ over \mathbb{R}^p by fitting different but simple model separately at each query point x_0 using only those observations close to the target point x_0 to fit the simple model. This localization is achieved via a weighting function kernel $K_\lambda(x_0, x_i)$, which assigns weight to x_i based on its distance from x_0 .

1. One-dimensional Kernel Smoothers

KNN uses unweighted average in the neighbourhood which caused the bumpy estimation. One can assign weights that die off smoothly within distance from the target point.

Kernel function

$$K_\lambda(x_0, x) = D\left(\frac{|x - x_0|}{h_\lambda(x_0)}\right)$$

Locally weighted regression (at each target point x_0):

$$\min_{\alpha(x_0), \beta(x_0)} = \sum_{i=1}^N K_\lambda(x_0, x_i) [y_i - \alpha(x_0) - \beta(x_0)x_i]^2$$

2. Selecting Width of the Kernel

- small window \rightarrow variance and small bias
- wide window the other way around

3. Kernel Density Estimation and Classification

Kernel density: $\hat{f}_X(x_0) = \frac{1}{N\lambda} \sum_{i=1}^N K_\lambda(x_0, x_i)$. A popular choice of K_λ is the Gaussian kernel $\phi(|x - x_0|/\lambda)$.

Suppose for a J class problem, we fit nonparametric density estimates $\hat{f}_j(X)$, $j = 1 \dots J$ in each class, and we know the class priors $\hat{\pi}_j$, then $\hat{P}_r(G = j|x_0) = \frac{\hat{\pi}_j \hat{f}_j(x_0)}{\sum_k \hat{\pi}_k \hat{f}_k(x_0)}$

3.1 Naive Bayes

Assumption: Given a class $G = j$, the features X_k in the p dimensional space are independent, i.e. $f_j(X) = \prod_{k=1}^p f_{jk}(X_k)$

- f_{jk} can be estimated using one dimensional kernel density estimates.

4. Radial Basis Functions and Kernels

Radial basis function treats kernel functions $K_\lambda(\xi, x)$ as basis functions for a smoothing spline estimation, which leads to

$$f(x) = \sum_{j=1}^M K_{\lambda_j}(\xi_j, x) \beta_j$$

Chapter 7: Model Assessment and Selection

1. Introduction

The performance of a learning method relates to its prediction capability on independent test data

2. Bias, Variance, and Model Complexity

2.1 Regression

- **Test error:** aka *generalization error*, is the prediction error over an independent test sample $\text{Err}_\tau = E[L(Y, \hat{f}(x)) | \tau]$. τ represents a fixed training set.
- **Expected prediction error:** or *expected test error* $\text{Err} = E[L(Y, \hat{f}(X))] = E[\text{Err}_\tau]$. Expectation averages over everything that is random, including the randomness in the training set that produced \hat{f} .
- **Training error:** is the average loss over the training sample $\text{err} = \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}(x_i))$

2.2 Classification

Loss function can be 0-1 loss or $-2 \times \log$ likelihood (aka deviance.)

2.3 Split of data

Training (fit the model), validation (model selection), and test (model assessment).

3. Bias-Variance Decomposition

The expected prediction error of a regression fit $\hat{f}(x)$ at input $X = x_0$

$$\begin{aligned}\text{Err}(x_0) &= E[(Y - \hat{f}(x_0))^2 | X = x_0] \\ &= \sigma_\epsilon^2 + [E\hat{f}(x_0) - f(x_0)]^2 + E[\hat{f}(x_0) - E\hat{f}(x_0)]^2 \\ &= \sigma_\epsilon^2 + \text{Bias}^2(\hat{f}(x_0)) + \text{Var}(\hat{f}(x_0))\end{aligned}$$

4. The Optimism of the Training Error Rate

5. Estimates of In-Sample Prediction Error

$C_p = \text{err} + 2 \frac{d}{N} \hat{\sigma}_\epsilon^2$, d is the number of inputs or basis functions

$$AIC = -\frac{2}{N} \log \text{lik} + 2 \frac{d}{N}$$

$AIC(\alpha) = \text{err}(\alpha) + 2 \frac{d(\alpha)}{N} \hat{\sigma}_\epsilon^2$ for models with tuning parameter α

6. Effective Number of Parameters

7. The Bayesian Approach and BIC

General form $\text{BIC} = -2 \log \text{lik} + (\log N)d$

10. Cross Validation

Goal: To estimate prediction error. It estimates the expected extra-sample error $\text{Err} = E[L(Y, \hat{f}(X))]$

10.1 K Fold CV

Split the data into K parts with equal size. For the k th part, we fit the model to the other $K - 1$ parts of the data and calculate the prediction error of the fitted model when predicting the k th part of the data. We do this for $k = 1, 2, \dots, K$ and combine the K estimates of the prediction errors.

$\parallel : \{1 \dots N\} \rightarrow \{1 \dots K\}$ represents the partition. $\hat{f}^{-k}(x)$ denoted the fitted value while excluding the set k . The cross validation estimate of the prediction error is

$$\text{CV}(\hat{f}) = \frac{1}{N} \sum_{i=1}^N L(y_i - \hat{f}^{-k}(x))$$

or

$$\text{CV}(\hat{f}) = \frac{1}{K} \sum_{k=1}^K \left[\frac{1}{n_k} \sum_{i \in \text{kth val set}} L(y_i - \hat{f}^{-k}(x)) \right]$$

LOCV: Low bias but high variance, also high computational burden.

One Standard Error Rule Choose the most parsimonious model whose error is no more than one standard error above the error of the best model.

Generalized Cross Validation A convenient approximation to leave one out CV, for linear fitting under squared error loss.

$$GCV(\hat{f}) = \frac{1}{N} \sum_i \left[\frac{y_i - \hat{f}(x_i)}{1 - \text{trace}(\mathbf{S})/N} \right]^2$$

10.2 Tips for CV

DON'T subset the data before CV, because the subset was chosen on the basis of all of the samples. Leaving samples out after the subset does not correctly mimic the application of the classifier.

In general, with a multistep modeling procedure, cross validation must be applied to the entire sequence of modeling steps.

11. Bootstrap

Chapter 8: Model Inference and Averaging

1. Introduction

- minimize sum of square for regression
- minimize cross entropy for classification

2. The Bootstrap and Maximum Likelihood Methods

2.1 A Smoothing Example

Denote the training data by $\mathbf{Z} = \{z_1, z_2, \dots, z_n\}$, with $z_i = (x_i, y_i)$, and $x_i \in \mathbb{R}^1$. If we use a B-spline basis expansion with 3 knots placed on the quartiles, we can represent the function by

$$E(Y|X = x) = \mu(x) = \sum_{j=1}^7 \beta_j h_j(x)$$

MSE

The MSE can be represented by $\hat{\beta} = (\mathbf{H}^T \mathbf{H})^{-1} \mathbf{H}^T \mathbf{y}$, where $\mathbf{H} \in \mathbb{R}^{N \times 7}$ with ij th elements $h_j(x_i)$. The estimated covariance matrix of $\hat{\beta}$ is $\text{Var}(\hat{\beta}) = (\mathbf{H}^T \mathbf{H})^{-1} \hat{\sigma}^2$, where $\hat{\sigma}^2 = \sum_{i=1}^N (y_i - \hat{\mu}(x_i))^2 / N$. The standard error of a prediction $\hat{\mu}(x) = h(x) \hat{\beta}$ is $\text{SE}[\hat{\mu}(x)] = [h(x)^T (\mathbf{H}^T \mathbf{H})^{-1} h(x)]^{1/2} \hat{\sigma}$. This can be used to construct the confidence interval.

Bootstrap Approach

Nonparametric bootstrap: Draw B datasets each of size $N = 50$ with replacement, where the sampling unit is $z_i = (x_i, y_i)$. Using each bootstrapped dataset \mathbf{Z}^* , we fit a cubic spline $\hat{\mu}^*(x)$. With large enough resamples, we can use percentiles to get the confidence band.

Parametric bootstrap: simulate new responses by adding Gaussian noise to the predicted values (assuming we assume a Gaussian error for the original model).

2.2 MLE

In general parametric bootstrap agree with MLE.

ADD MLE FURTHER

3. Bayesian Methods

Baysian specifies a sampling model $\text{Pr}(\mathbf{Z}|\theta)$, and a prior distribution for the parameter $\text{Pr}(\theta)$ representing our knowledge about θ . The goal is to compute the posterior distribution

$$\text{Pr}(\theta|\mathbf{Z}) = \frac{\text{Pr}(\mathbf{Z}|\theta) \cdot \text{Pr}(\theta)}{\int \text{Pr}(\mathbf{Z}|\theta) \cdot \text{Pr}(\theta) d\theta}$$

.

A future observation can be predicted via the predictive distribution

$$\text{Pr}(z^{new}|\mathbf{Z}) = \int \text{Pr}(z^{new}|\theta) \text{Pr}(\theta|\mathbf{Z}) d\theta$$

4. Bootstrap v.s Bayesian Inference

5. The EM Algorithm

5.1 Two-Component Mixture Model

5.2 EM in General

5.3. EM as a Maximization - Maximzation Procedure

6. MCMC for Sampling from the Posterior

7. Bagging

Bootstrap mean is approximately a posterior average,

Consider training data $\mathbf{Z} = \{(x_1, y_1), \dots, (x_N, y_N)\}$ and prediction $\hat{f}(x)$ at input x . Bagging (bootstrap aggregation) averages this prediction over a collection of bootstrap samples, thereby reducing its variance. For each bootstrap sample $\mathbf{Z}^{*b}, b = 1 \dots B$, we fit the model and get the prediction $\hat{f}^{*b}(x)$. The bagging estimate is defined by

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

For a classification problem with K responses, and the predicted classes $\hat{G}(x)$. For each single prediction, it provides an indicator vector with value a single one and $K-1$ zeroes. The the bagged estimation is a K -vector $[p_1(x), p_2(x), \dots, p_K(x)]$, with $p_k(x)$ equal to the proportion of predicted class k at x .

8. Model Averaging and Stacking

Chapter 9: Additive Models, Trees, and Related Methods

1. Generalized Additive Model

Definition

$$E(Y|X_1, \dots, X_p) = \alpha + f_1(X_1) + f_2(X_2) + \dots + f_p(X_p)$$

where X_k are predictors, and $f_k(\cdot)$ are nonparametric smooth functions.

1.1 Fitting Additive Models

In order to get unique solution, the standard convention is to assume that $\sum_{i=1}^N f_j(x_{ij}) = 0 \quad \forall j$. The $\hat{\alpha} = \text{ave}(y_i)$.

Algorithm 9.1 *The Backfitting Algorithm for Additive Models.*

1. Initialize: $\hat{\alpha} = \frac{1}{N} \sum_1^N y_i$, $\hat{f}_j \equiv 0, \forall i, j$.

2. Cycle: $j = 1, 2, \dots, p, \dots, 1, 2, \dots, p, \dots$,

$$\begin{aligned} \hat{f}_j &\leftarrow \mathcal{S}_j \left[\{y_i - \hat{\alpha} - \sum_{k \neq j} \hat{f}_k(x_{ik})\}_1^N \right], \\ \hat{f}_j &\leftarrow \hat{f}_j - \frac{1}{N} \sum_{i=1}^N \hat{f}_j(x_{ij}). \end{aligned}$$

until the functions \hat{f}_j change less than a prespecified threshold.

Figure 7: The Backfitting Algorithm for Additive Models

1.2 Additive Logistic Regression

2. Tree-Based Methods

2.1 Background

Recursive binary partitions

2.2 Regression Trees

Data (x_i, y_i) with $i = 1 \dots N$, and $x_i = c(x_{i1} \dots x_{ip})$, i.e. p features.

Suppose we have a partition into M regions, the response can be modeled as a constant c_m in each region, $f(x) = \sum_{m=1}^M c_m I(x \in R_m)$. LSE gives $\hat{c}_m = \text{ave}(y_i | x_i \in R_m)$.

Algorithm 9.2 *Local Scoring Algorithm for the Additive Logistic Regression Model.*

1. Compute starting values: $\hat{\alpha} = \log[\bar{y}/(1 - \bar{y})]$, where $\bar{y} = \text{ave}(y_i)$, the sample proportion of ones, and set $\hat{f}_j \equiv 0 \forall j$.
2. Define $\hat{\eta}_i = \hat{\alpha} + \sum_j \hat{f}_j(x_{ij})$ and $\hat{p}_i = 1/[1 + \exp(-\hat{\eta}_i)]$.

Iterate:

- (a) Construct the working target variable

$$z_i = \hat{\eta}_i + \frac{(y_i - \hat{p}_i)}{\hat{p}_i(1 - \hat{p}_i)}.$$

- (b) Construct weights $w_i = \hat{p}_i(1 - \hat{p}_i)$
- (c) Fit an additive model to the targets z_i with weights w_i , using a weighted backfitting algorithm. This gives new estimates $\hat{\alpha}, \hat{f}_j, \forall j$

3. Continue step 2. until the change in the functions falls below a pre-specified threshold.
-

Figure 8: Local Scoring Algorithm for the Additive Logistic Regression Model

We can use a greedy algorithm to start with all of the data, consider a splitting variable j and split point s , and find the optimal variable j and s .

Cost Complexity Pruning: To grow a large tree T_0 and stop when a minimum node size is reached. For any tree T which is a subtree of T_0 , the cost complexity function is

$$C_\alpha(T) = \sum_{m=1}^{|T|} N_m Q_m(T) + \alpha |T|$$

where N_m is number of observations falling into terminal node variable, $Q_m()$ is the variance, and α is a tuning parameter, which controls the tree size and goodness of fit.

2.3 Classification Tree

In a node m representing a region R_m with N_m observations, let $\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \in R_m} I(y_i = k)$ be the proportion of class k observations in node m . We classify the observations in node m to class $k(m) = \operatorname{argmax}_k \hat{p}_{mk}$, i.e. the majority class in node m . Measures of node impurity include

- Misclassification error: $1 - \hat{p}_{mk(m)}$
- Gini index: $\sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$
- Cross entropy or deviance: $-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$

Gini and CE are more sensitive to changes in the node probabilities than error rate, so should be used.

2.4 Practical Issues

- The partitioning algorithm tends to favor categorical predictors with many levels q ; the number of partitions grows exponentially in q , and the more choices we have, the more likely we can find a good one for the data at hand, which leads to severe over-fitting if q is large.
- Consequences of misclassifying observations may have different prices. To account for this, we can consider a $K \times K$ loss matrix \mathbf{K} with $L_{kk'}$ being the loss for classifying a class k observation as class k' . Diagonal is typically 1. Gini can be then modified to $\sum_{k \neq k'} L_{kk'} \hat{p}_{mk} \hat{p}_{mk'}$.
- Missing categorical variable: add missing category. Or to construct surrogate variable
- Multi-way partition is not recommended.
- Trees are difficult to model additive structure.

3. PRIM: Bump Hunting

4. MARS; Multivariate Adaptive Regression Splines

TBC

5. Hierarchical Mixtures of Experts

TBC

Algorithm 9.3 *Patient Rule Induction Method.*

1. Start with all of the training data, and a maximal box containing all of the data.
 2. Consider shrinking the box by compressing one face, so as to peel off the proportion α of observations having either the highest values of a predictor X_j , or the lowest. Choose the peeling that produces the highest response mean in the remaining box. (Typically $\alpha = 0.05$ or 0.10 .)
 3. Repeat step 2 until some minimal number of observations (say 10) remain in the box.
 4. Expand the box along any face, as long as the resulting box mean increases.
 5. Steps 1–4 give a sequence of boxes, with different numbers of observations in each box. Use cross-validation to choose a member of the sequence. Call the box B_1 .
 6. Remove the data in box B_1 from the dataset and repeat steps 2–5 to obtain a second box, and continue to get as many boxes as desired.
-

Figure 9: PRIM

6. Missing Data

Missing at Random $Pr(R|Z, \theta) = Pr(R|Z_{obs}, \theta)$

Missing completely at random $Pr(R|Z, \theta) = Pr(R|\theta)$

Chapter 10: Boosting and Additive Trees

1. Introduction

Chapter 11: Neural Networks

1. Introduction

Chapter 12: Support Vector Machines and Flexible Discriminants

1. Introduction

Chapter 13: Prototype Methods and Nearest-Neighbors

1. Introduction

Chapter 14: Unsupervised Learning

1. Introduction

Chapter 15: Random Forests

1. Introduction

Chapter 16: Ensemble Learning

1. Introduction

Chapter 17: Undirected Graphical Models

1. Introduction

Chapter 18: High-Dimensional Problems: $p \gg N$

1. Introduction